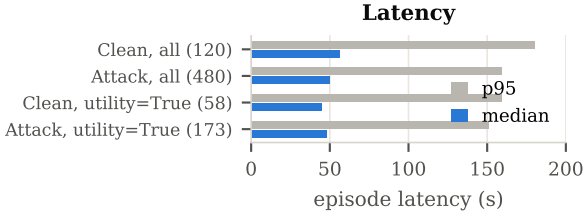
Latency
Clean, all (120)
Attack, all (480)
Clean, utility=True (58)
Attack, utility=True (173)
p95
median
0
50
100
150
200
episode latency (s)

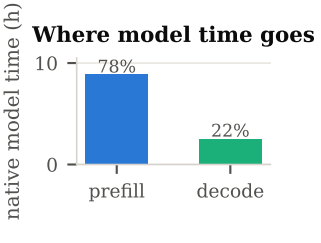
Where model time goes
native model time (h)
10
0
78%
22%
prefill
decode

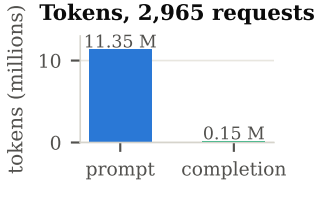
Tokens, 2,965 requests
tokens (millions)
11.35 M
10
0
0.15 M
prompt
completion