Latency
Clean, all (120)
Attack, all (480)
Clean, utility=True (58)
Attack, utility=True (173)
0
50
100
150
200
episode latency (s)
Where model time goes
native model time (h)
10.0
7.5
5.0
2.5
0.0
78.8%
21.2%
prefill
decode
Tokens, 2,927 responses
tokens (millions)
10
5
0
11.12 M
0.14 M
prompt
completion